


Breadth-First Search (BFS)


Intuition: BFS(vertex v)

To start, all vertices are *unmarked*.

- *Start* at v . *Visit* v and mark as *visited*.
- *Visit* every *unmarked* neighbour u_i of v and mark each u_i as *visited*.
- *Mark* v *finished*.
- *Recurse* on each vertex marked as *visited* in the order they were visited.



BFS of an undirected Graph



Q: What *information* about the graph can a *BFS* be used to find?

-
-

Q: What does the *BFS* construct?

Q: What is an appropriate *ADT* to *implement* a *BFS* given an *adjacency list* representation of a graph?

which has the operations:

-
-
-

Q: What *information* will we need to store along the way?

-
-
-

The BFS Algorithm

We will use $p[v]$ to represent the *predecessor* of v and $d[v]$ to represent the *number of edges* from v (i.e., the *distance* from v).

```
BFS (G= (V, E) , v)
  for all vertices u in V
    color[u] := black
    d[u] := infinity;      \\ we use infinity to denote
    p[u] := NIL;          \\ "not connected"
  end for
  initialize an empty queue Q;
  color[v] := green;
  d[v] := 0;
  p[v] := NIL;
  ENQUEUE (Q, v);
  while not ISEMPTY(Q) do
    u := DEQUEUE (Q);
    for each edge (u, w) in E do
      if (color[w] == black) then
        color[w] := green;
        d[w] := d[u] + 1;
        p[w] := u;
        ENQUEUE (Q, w);
      end if
    end for
    color[u] := red;
  end while
END BFS
```

Complexity of BFS(G, v)

Q: How many times is each node **ENQUEUE**ed?

Therefore, the *adjacency list of each node* is examined at most once, so that the total running time of **BFS** is

or **linear** in the size of the adjacency list.

NOTES:

- BFS will visit only those vertices that are *reachable* from v .
- If the graph is *connected (in the undirected case)* or *strongly-connected (in the directed case)*, then this will be all the vertices.
- If not, then we may have to call BFS on more than one start vertex in order to see the whole graph.

Q: Prove that $d[u]$ really does represent the *length* of the *shortest path* (in terms of number of edges) from v to u .


Depth-First Search

Intuition: DFS(G, v)

All *vertices* and *edges* start out *unmarked*.

- Walk as *far* as possible *away* from v visiting vertices
- If the *current vertex* has not been *visited*,
 - Mark as *visited* and the *edge* that is traversed as a **DFS edge**.
- Otherwise, if the *vertex* has been *visited*,
 - mark the *traversed edge* as a **back-edge**, back up to the previous vertex
- When the *current vertex* has only *visited neighbours left* mark as *finished (white)*.
- *Backtrack* to the *first vertex* that is not *finished*.
- Continue.

Example.



Just like BFS, *DFS* constructs a spanning-tree and gives *connected component* information.

Q: Does it find the *shortest path* between v and all other *vertices*?

Implementing a DFS

Q: Which *ADT* would be the most helpful for *implementing DFS* given an *adjacency list* representation of G ?

with the operations

-
-
-

Q: What *additional data* (for each vertex) should we keep in order to easily determine *whether an edge* is a *cross-edge* or a *DFS-edge*?

- $d[v]$ will indicate the
- $f[v]$ will indicate the

Algorithm DFS(G,s)

```
DFS (G=(V,E),s)
  for all vertices v in V
    color[v] := black;
    d[v] := infinity;
    f[v] := infinity;
    p[v] := NIL;
  end for
  initialize an empty stack S;
  color[s] := green; d[s] := 0; p[s] := NIL;
  time := 0;
  PUSH(S, (s,NIL));
  for each edge (s,v) in E do
    PUSH(S, (s,v));
  end for
  while not ISEMPY(S) do
    (u,v) := POP(S);
    if (v == NIL) then           // Done with u
      time := time + 1;
      f[u] := time;
      color[u] := white;
    end if
    else if (color[v] == black) then
      color[v] := green;
      time := time + 1;
      d[v] := time;
      p[v] := u;
      PUSH(S, (v,NIL));         // Marks the end of v's neighbors
      for each edge (v,w) in E do
        PUSH(S, (v,w));
      end for
    (*) end if
  end while
END DFS
```


Complexity of DFS(G, s)

Q: How many times does *DFS* visit the *neighbours* of a *node*?

-
- Therefore, the *adjacency list* of each *vertex* is visited at most once.
- So the *total running time* is just like for BFS, $\Theta(n + m)$ i.e., *linear* in the *size* of the *adjacency list*.

Note that the *gold* edges, or the *DFS edges* form a *tree* called the *DFS-tree*.

Q: Is the *DFS tree* unique for a given graph G starting at s ?

For certain applications, we need to *distinguish* between different types of *edges* in E :

We can specify edges on *DFS-tree* according to how they are *traversed* during the search.

- **Tree-Edges** are the edges in the *DFS tree*.
- **Back-Edges** are edges from a vertex u to an *ancestor* of u in the DFS tree.
- **Forward-Edges**^{*} are edges from a vertex u to a *descendent* of u in the DFS tree.
- **Cross-Edges**^{*} are all the *other edges* that are *not* part of the *DFS tree* (from a vertex u to another vertex v that is neither an ancestor nor a descendent of u in the DFS tree).

Q: Which *variable* facilitates *distinguishing* between these edges?

Q: How can a *DFS* be used to *determine* whether a graph G has any *cycles*?


(Note: A *cycle* is a path from a vertex u to itself.)

Q: How can we detect *cross-edges* during a *DFS*?


Minimum Cost Spanning Trees (MCSTs)

- Let $G = (V, E)$ be a *connected, undirected* graph with *edge weights* $w(e)$ for each edge $e \in E$.
- A *tree* is a subset of edges $A \subset E$ such that A is *connected* and *does not contain a cycle*.

Thick edges are in A , the thin ones are not. The first one is a tree and the other two are not.



A *spanning tree* is a tree A such that *every vertex* $v \in V$ is an *endpoint* of at least *one edge* in A .



Q: How *many edges* must any *spanning tree* contain?

How would you prove this is true?


Q: What algorithms have we seen to *construct a spanning tree*?

A *minimum cost spanning tree* (**MCST**) is a spanning tree *A* such that the *sum of the weights is minimized* for all possible spanning trees *B*.


Formally:

$$w(A) = \sum_{e \in A} w(e) \leq w(B)$$



for all other spanning trees *B*.



Minimum Cost Spanning Tree



Spanning tree, but not minimum cost



Two different MCSTs on the same graph

Q: What is an *example* of an *application* in which you would want to be able to find a *MCST*?

A:

We will look at two algorithms for constructing MCSTs. The first is *Prim's Algorithm*.

Prim's Algorithm

Prim's algorithm uses a *Priority Queue ADT*.

A *priority queue* is just like a queue except that *every item in the queue* has a *priority* (usually just a number).

More formally, a *priority queue* consists of

- a set of elements
- each element has a priority
- The operations:
 - `ENQUEUE (x, p) :`
 - `ISEMPTY () :`
 - `EXTRACT_MIN () :`